Graph Algorithms

Graphs

- \triangleright A graph G = (V, E)
 - \blacksquare V = set of vertices
 - \blacksquare E = set of edges = subset of V × V
 - Thus $|E| = O(|V|^2)$

Graph Variations

- > Variations:
 - A connected graph has a path from every vertex to every other
 - In an *undirected graph:*
 - Edge (u,v) = edge(v,u)
 - No self-loops
 - In a *directed* graph:
 - Edge (u,v) goes from vertex u to vertex v, notated $u\rightarrow v$

Graph Variations

- More variations:
 - A *weighted graph* associates weights with either the edges or the vertices
 - E.g., a road map: edges might be weighted w/ distance
 - A *multigraph* allows multiple edges between the same vertices
 - E.g., the call graph in a program (a function can get called from multiple points in another function)

Graphs

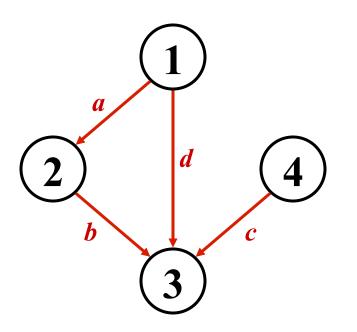
- ➤ We will typically express running times in terms of |E| and |V| (often dropping the |'s)
 - If $|E| \approx |V|^2$ the graph is *dense*
 - If $|E| \approx |V|$ the graph is *sparse*
- If you know you are dealing with dense or sparse graphs, different data structures may make sense

Representing Graphs

ightharpoonup Assume V = {1, 2, ..., n}

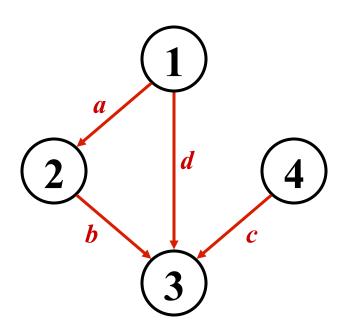
- An *adjacency matrix* represents the graph as a $n \times n$ matrix A:
 - A[i,j] = 1 if edge $(i,j) \in E$ (or weight of edge) = 0 if edge $(i,j) \notin E$

> Example:



A	1	2	3	4
1				
2				
3			??	
4				

> Example:



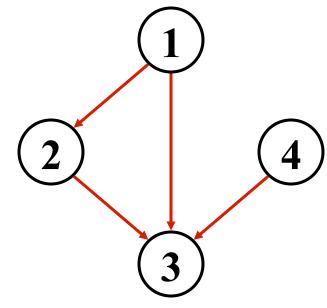
A	1	2	3	4
1	0	1	1	0
2	0	0	1	0
3	0	0	0	0
4	0	0	1	0

- ➤ How much storage does the adjacency matrix require?
- \rightarrow A: O(V²)
- What is the minimum amount of storage needed by an adjacency matrix representation of an undirected graph with 4 vertices?
- > A: 6 bits
 - Undirected graph → matrix is symmetric
 - No self-loops → don't need diagonal

- The adjacency matrix is a dense representation
 - Usually too much storage for large graphs
 - But can be very efficient for small graphs
- Most large interesting graphs are sparse
 - E.g., planar graphs, in which no edges cross, have |E| = O(|V|) by Euler's formula
 - For this reason the *adjacency list* is often a more appropriate respresentation

Graphs: Adjacency List

- Adjacency list: for each vertex $v \in V$, store a list of vertices adjacent to v
- > Example:
 - \blacksquare Adj[1] = {2,3}
 - $Adj[2] = {3}$
 - \blacksquare Adj[3] = {}
 - \blacksquare Adj[4] = {3}
- Variation: can also keep a list of edges coming *into* vertex



Graphs: Adjacency List

- ➤ How much storage is required?
 - The *degree* of a vertex v = # incident edges
 - Directed graphs have in-degree, out-degree
 - For directed graphs, # of items in adjacency lists is Σ out-degree(v) = |E| takes $\Theta(V + E)$ storage
 - For undirected graphs, # items in adj lists is Σ degree(v) = 2 |E| also $\Theta(V + E)$ storage
- > So: Adjacency lists take O(V+E) storage

Graph Searching

- ➤ Given: a graph G = (V, E), directed or undirected
- Goal: methodically explore every vertex and every edge
- Ultimately: build a tree on the graph
 - Pick a vertex as the root
 - Choose certain edges to produce a tree
 - Note: might also build a *forest* if graph is not connected

Breadth-First Search

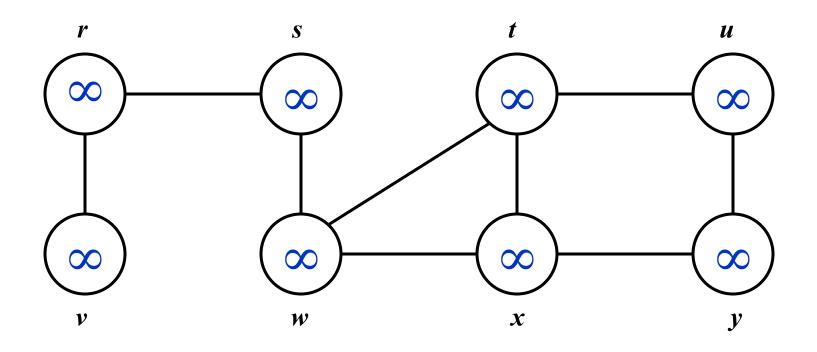
- > "Explore" a graph, turning it into a tree
 - One vertex at a time
 - Expand frontier of explored vertices across the breadth of the frontier
- Builds a tree over the graph
 - Pick a *source vertex* to be the root
 - Find ("discover") its children, then their children, etc.

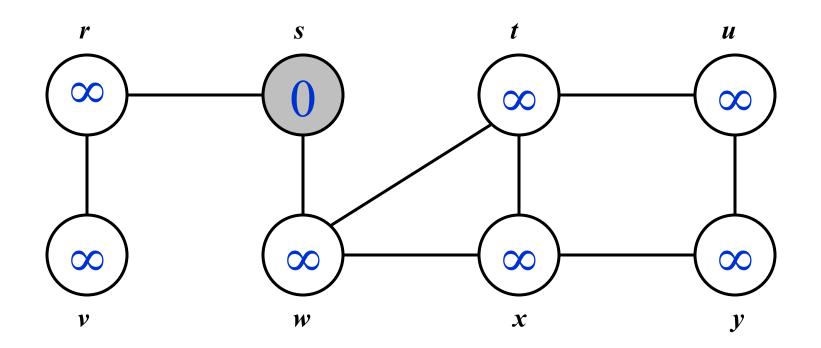
Breadth-First Search

- ➤ Again will associate vertex "colors" to guide the algorithm
 - White vertices have not been discovered
 - All vertices start out white
 - Grey vertices are discovered but not fully explored
 - They may be adjacent to white vertices
 - Black vertices are discovered and fully explored
 - They are adjacent only to black and gray vertices
- Explore vertices by scanning adjacency list of grey vertices

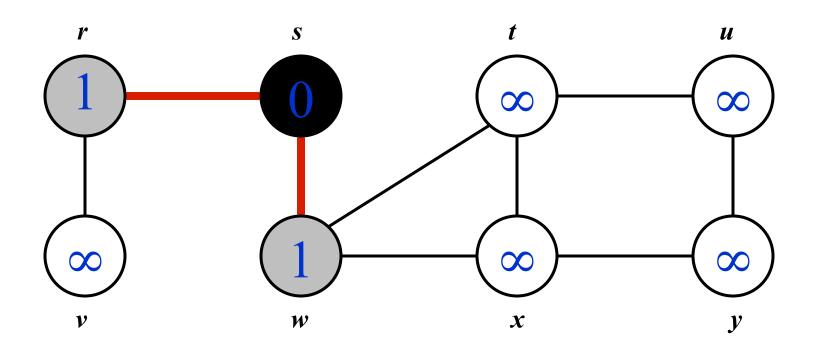
Breadth-First Search

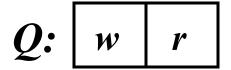
```
BFS(G, s) {
    initialize vertices;
    Q = \{s\}; // Q is a queue (duh); initialize to s
    while (Q not empty) {
        u = RemoveTop(Q);
        for each v \in u->adj {
            if (v->color == WHITE)
                v->color = GREY;
                v->d = u->d + 1:
                v->p = u;
                Enqueue(Q, v);
        u->color = BLACK;
```

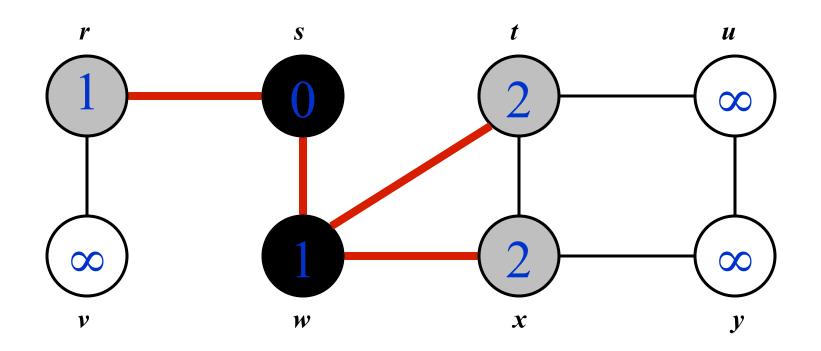




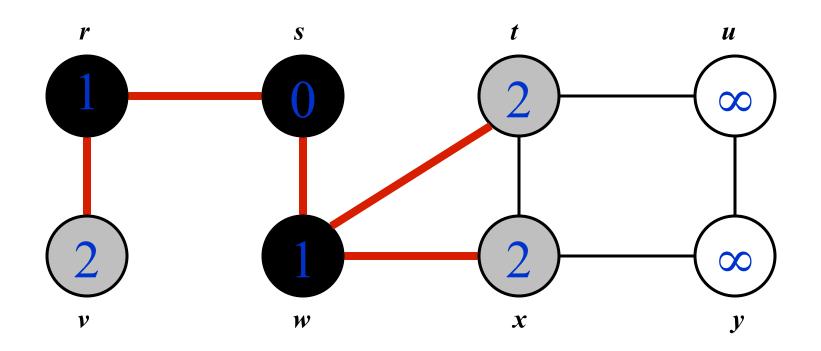
Q: s



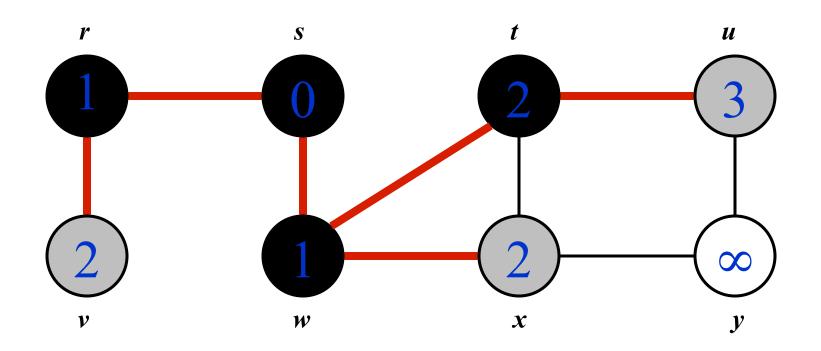


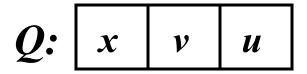


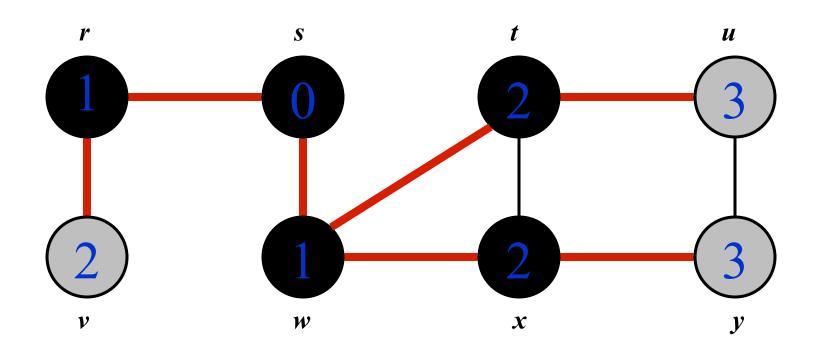
$$Q: r \mid t \mid x$$

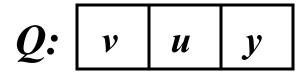


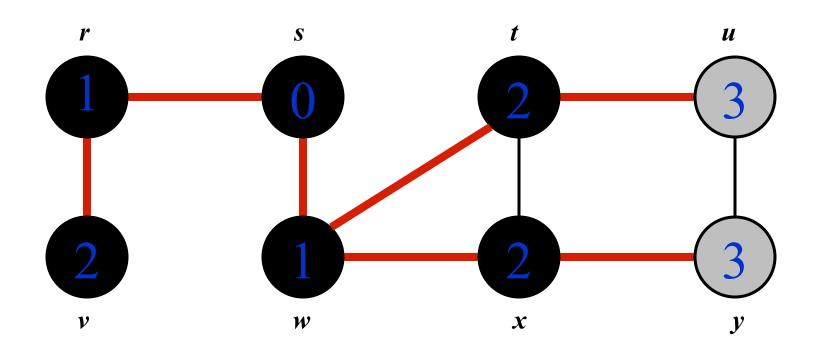
 $Q: \begin{array}{|c|c|c|c|c|} \hline t & x & v \\ \hline \end{array}$



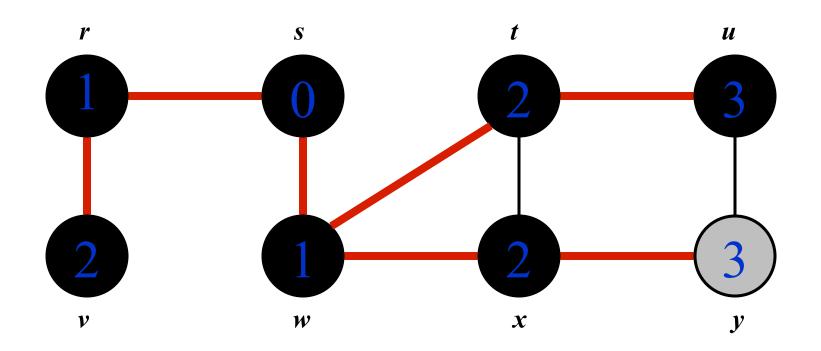




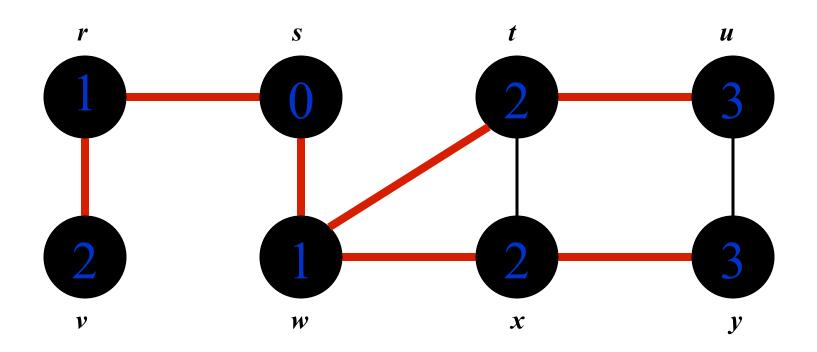




Q: | u | y |



Q: | y |



Q: Ø

BFS: The Code Again

```
BFS(G, s) {
       initialize vertices; \longleftarrow Touch every vertex: O(V)
      Q = \{s\};
      while (Q not empty) {
           u = RemoveTop(Q); \leftarrow u = every vertex, but only once
           for each v \in u-adj {
                                                          (Whv?)
               if (v->color == WHITE)
So v = every \ vertex \ v -> color = GREY;
that appears in v->d = u->d + 1;
some other vert's v->p = u;
                   Enqueue(Q, v);
adjacency list
                                   What will be the running time?
           u->color = BLACK;
                                   Total running time: O(V+E)
       }
```

BFS: The Code Again

```
BFS(G, s) {
    initialize vertices;
    Q = \{s\};
    while (Q not empty) {
        u = RemoveTop(Q);
        for each v \in u->adj {
             if (v->color == WHITE)
                 v->color = GREY;
                 v->d = u->d + 1:
                 v->p = u;
                 Enqueue(Q, v);
                                 What will be the storage cost
                                 in addition to storing the tree?
        u->color = BLACK;
                                 Total space used:
    }
                                 O(max(degree(v))) = O(E)
```

Breadth-First Search: Properties

- ➤ BFS calculates the *shortest-path distance* to the source node
 - Shortest-path distance $\delta(s,v)$ = minimum number of edges from s to v, or ∞ if v not reachable from s
 - Proof given in the book
- ➤ BFS builds *breadth-first tree*, in which paths to root represent shortest paths in G
 - Thus can use BFS to calculate shortest path from one vertex to another in O(V+E) time

Depth-First Search

- > Depth-first search is another strategy for exploring a graph
 - Explore "deeper" in the graph whenever possible
 - Edges are explored out of the most recently discovered vertex v that still has unexplored edges
 - When all of v's edges have been explored, backtrack to the vertex from which v was discovered

Depth-First Search

- Vertices initially colored white
- > Then colored gray when discovered
- > Then black when finished

```
DFS(G)
   for each vertex u \in G->V
      u->color = WHITE;
   time = 0;
   for each vertex u \in G->V
      if (u->color == WHITE)
         DFS Visit(u);
```

```
DFS Visit(u)
   u->color = GREY;
   time = time+1;
   u->d = time;
   for each v \in u-\lambda dj[]
       if (v->color == WHITE)
          DFS Visit(v);
   u->color = BLACK;
   time = time+1;
   u->f = time;
```

```
DFS(G)
   for each vertex u \in G->V
      u->color = WHITE;
   time = 0;
   for each vertex u \in G->V
      if (u->color == WHITE)
         DFS Visit(u);
```

```
DFS Visit(u)
   u->color = GREY;
   time = time+1;
   u->d = time;
   for each v \in u-\lambda dj[]
       if (v->color == WHITE)
          DFS Visit(v);
   u->color = BLACK;
   time = time+1;
   u->f = time;
```

```
DFS(G)
   for each vertex u \in G->V
      u->color = WHITE;
   time = 0;
   for each vertex u \in G->V
      if (u->color == WHITE)
         DFS Visit(u);
```

```
DFS Visit(u)
   u->color = GREY;
   time = time+1;
   u->d = time;
   for each v \in u-\lambda dj[]
       if (v->color == WHITE)
          DFS Visit(v);
   u->color = BLACK;
   time = time+1;
   u->f = time;
```

Running time: $O(n^2)$ because call DFS_Visit on each vertex, and the loop over Adj[] can run as many as |V| times

```
DFS(G)
   for each vertex u \in G->V
      u->color = WHITE;
   time = 0;
   for each vertex u \in G->V
      if (u->color == WHITE)
         DFS Visit(u);
```

```
DFS Visit(u)
   u->color = GREY;
   time = time+1;
   u->d = time;
   for each v \in u-\lambda dj[]
       if (v->color == WHITE)
          DFS Visit(v);
   u->color = BLACK;
   time = time+1;
   u->f = time;
```

BUT, there is actually a tighter bound.

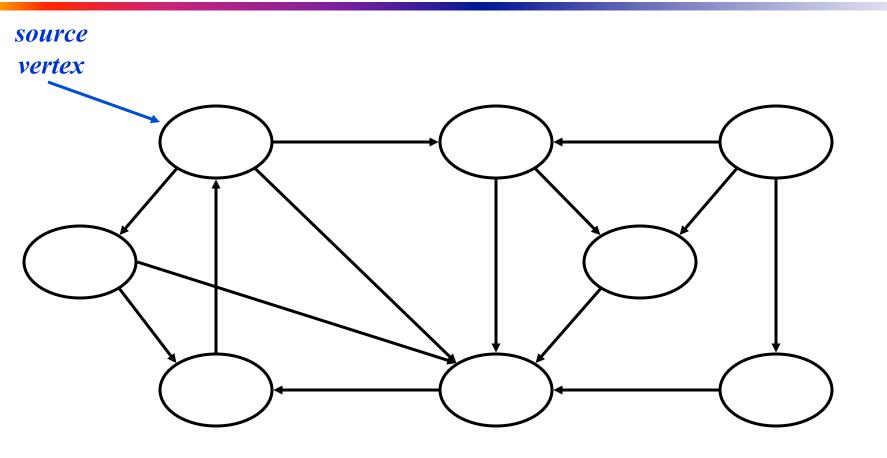
How many times will DFS_Visit() actually be called?

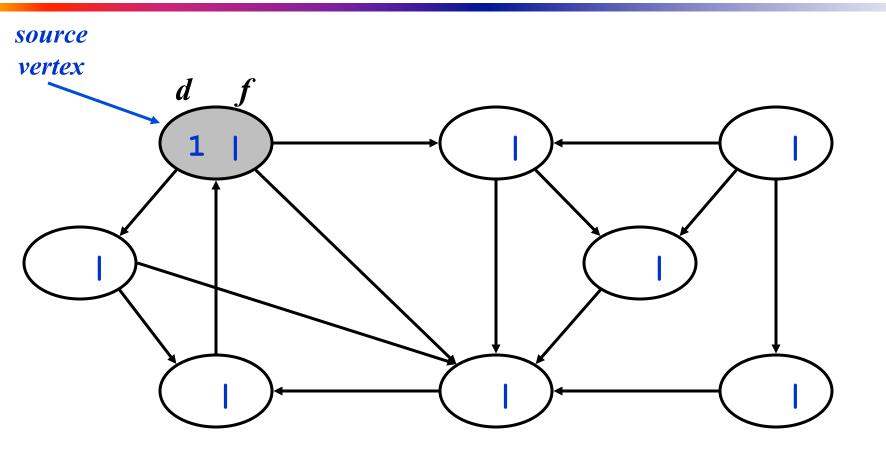
```
DFS(G)
   for each vertex u \in G->V
      u->color = WHITE;
   time = 0;
   for each vertex u \in G->V
      if (u->color == WHITE)
         DFS Visit(u);
```

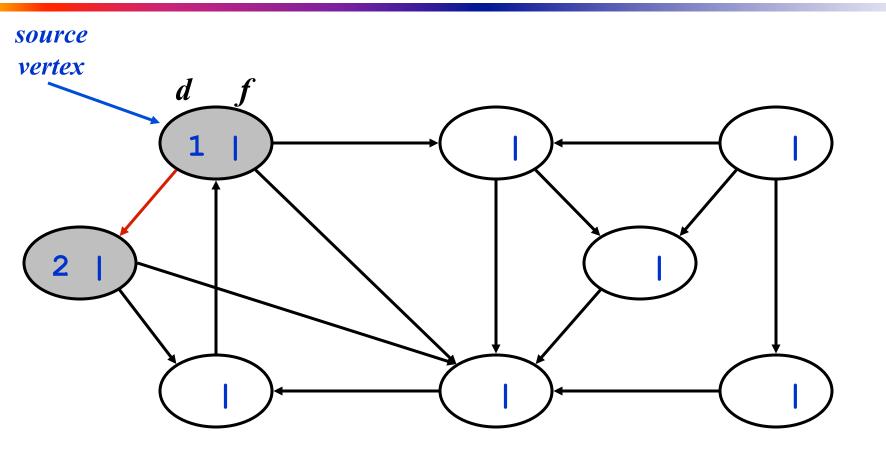
```
DFS Visit(u)
  u->color = GREY;
   time = time+1;
  u->d = time;
   for each v \in u-Adj[]
      if (v->color == WHITE)
          DFS Visit(v);
   u->color = BLACK;
   time = time+1;
   u->f = time;
```

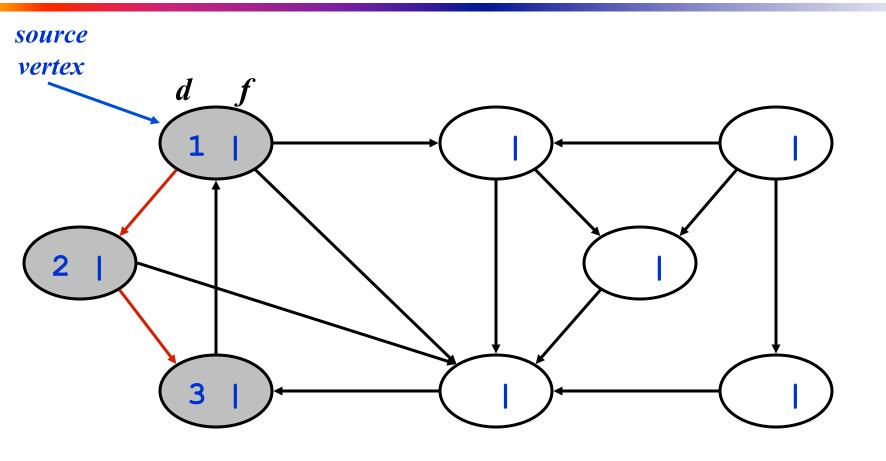
Depth-First Sort Analysis

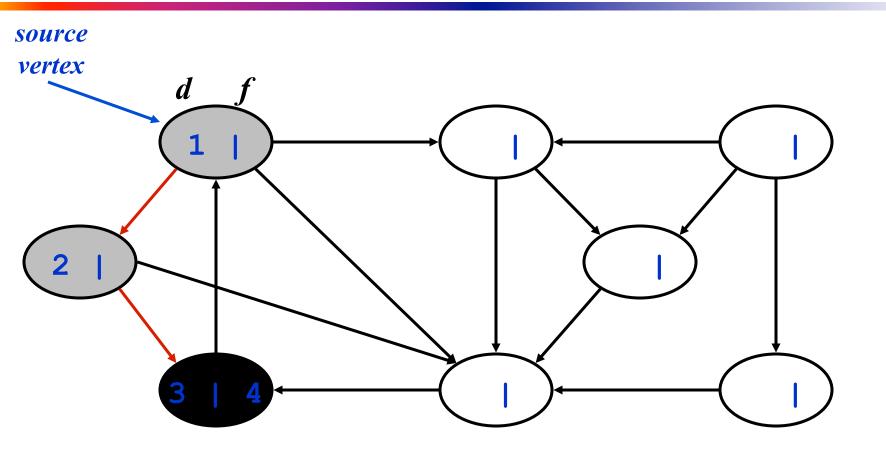
- This running time argument is an informal example of *amortized analysis*
 - "Charge" the exploration of edge to the edge:
 - Each loop in DFS_Visit can be attributed to an edge in the graph
 - Runs once/edge if directed graph, twice if undirected
 - Thus loop will run in O(E) time, algorithm O(V+E)
 - ◆ Considered linear for graph, b/c adj list requires O(V+E) storage
 - Important to be comfortable with this kind of reasoning and analysis

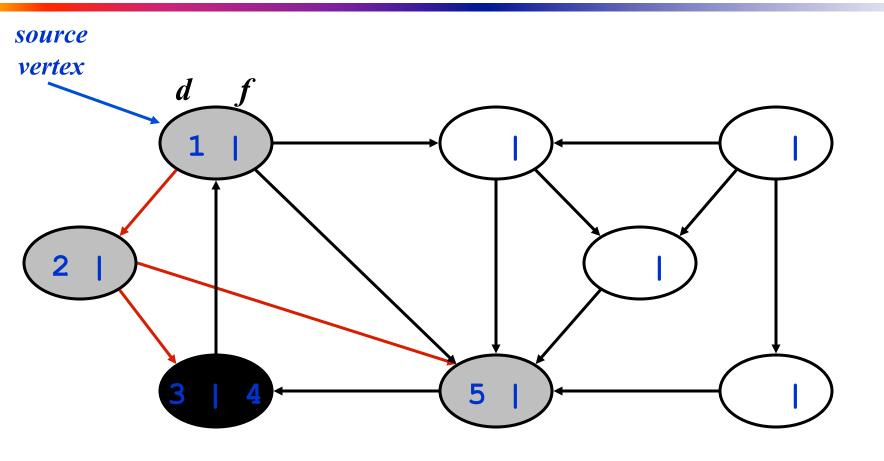


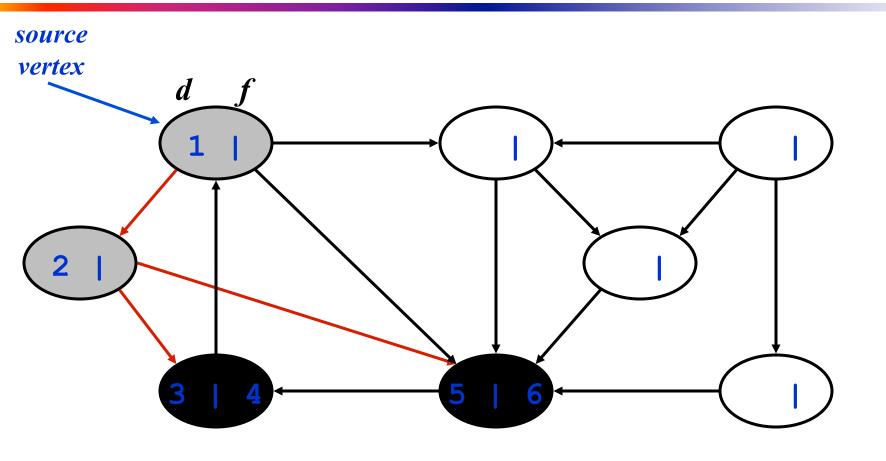


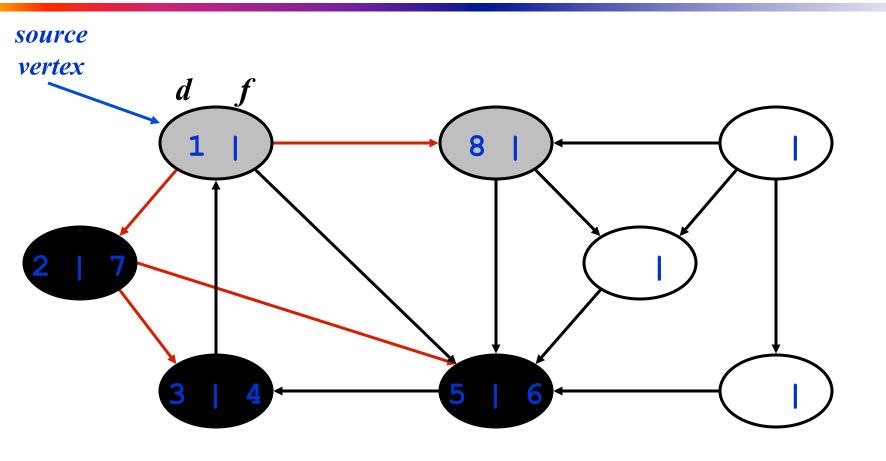


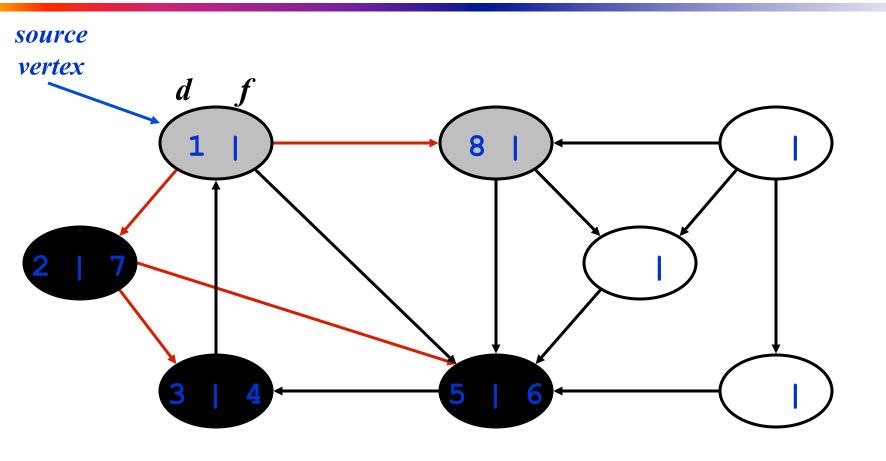


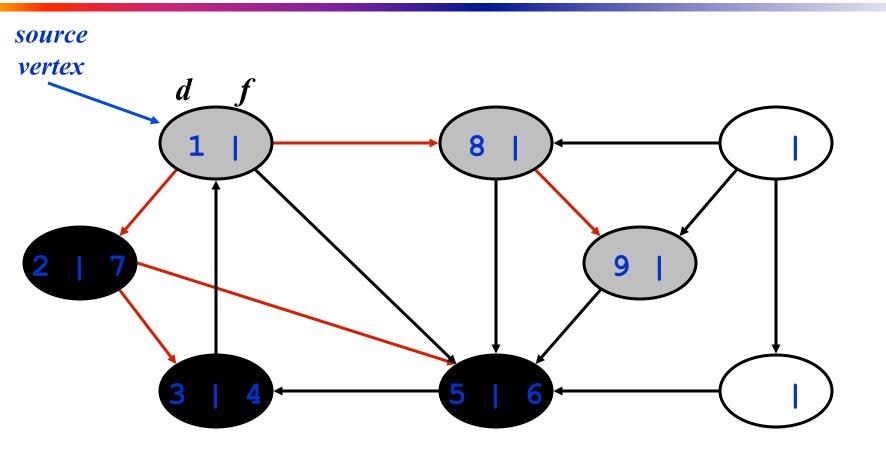


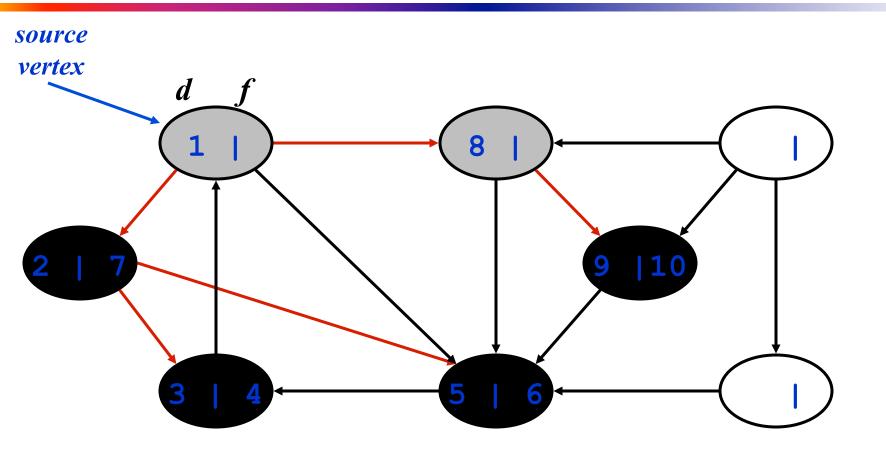


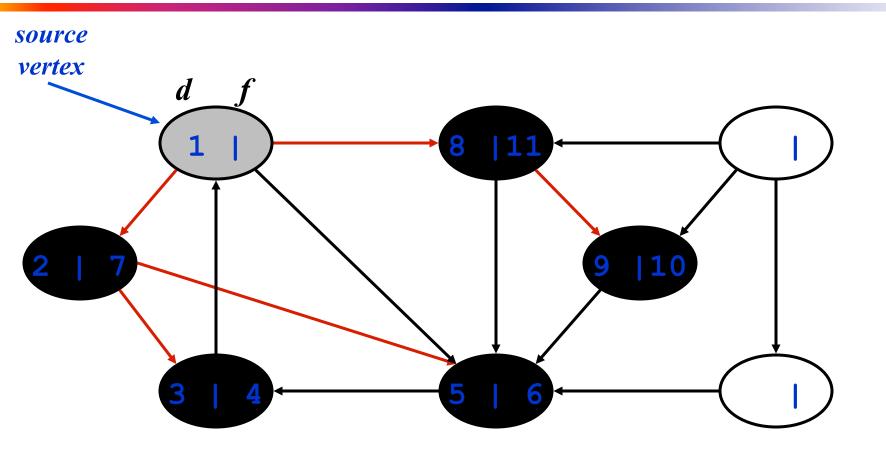


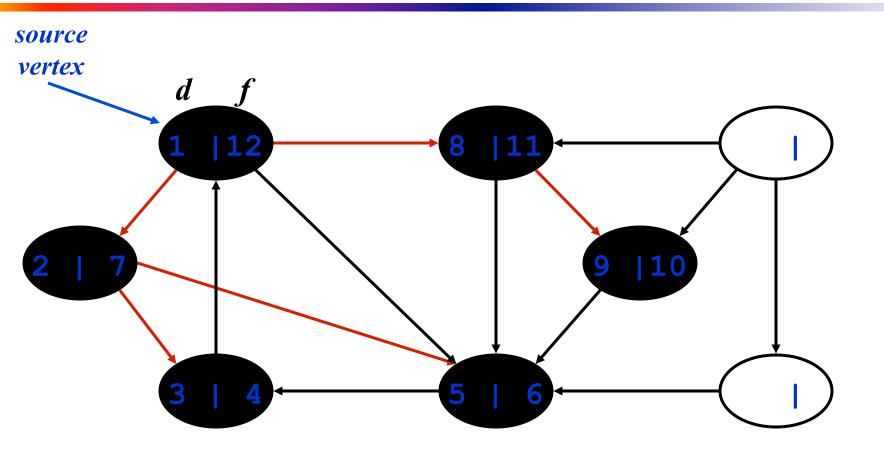


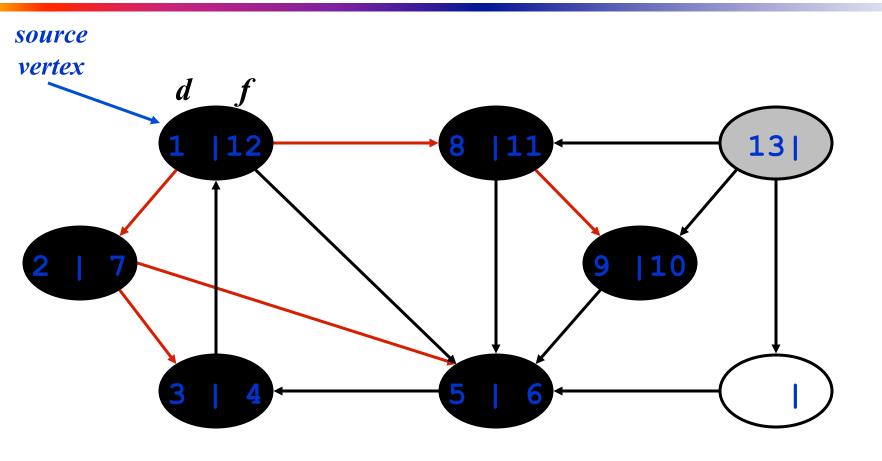


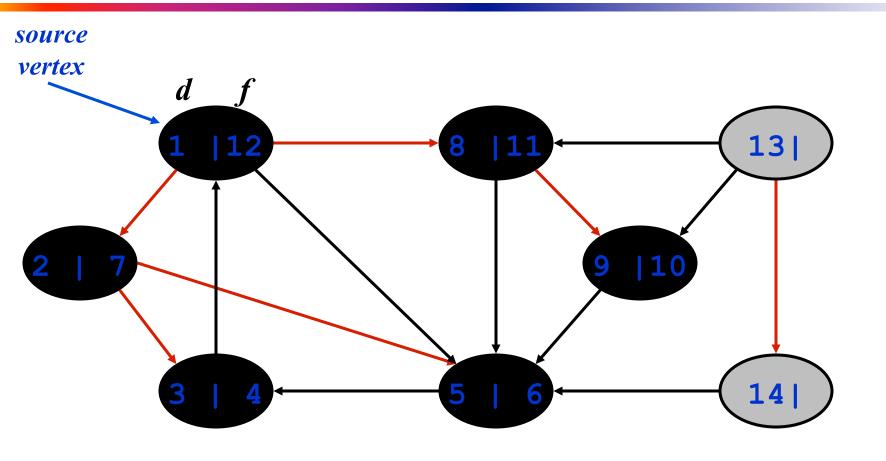


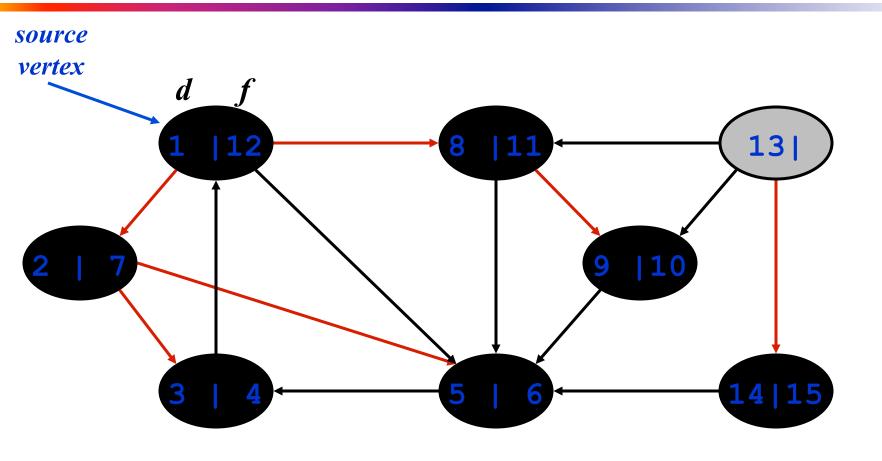


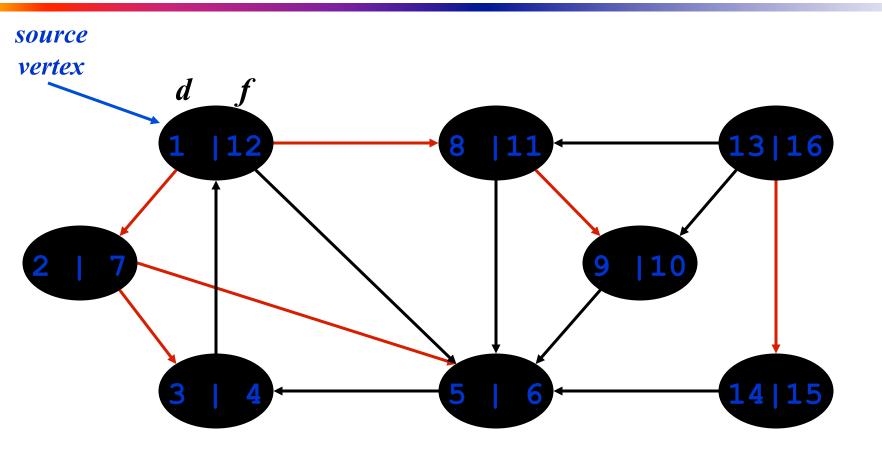






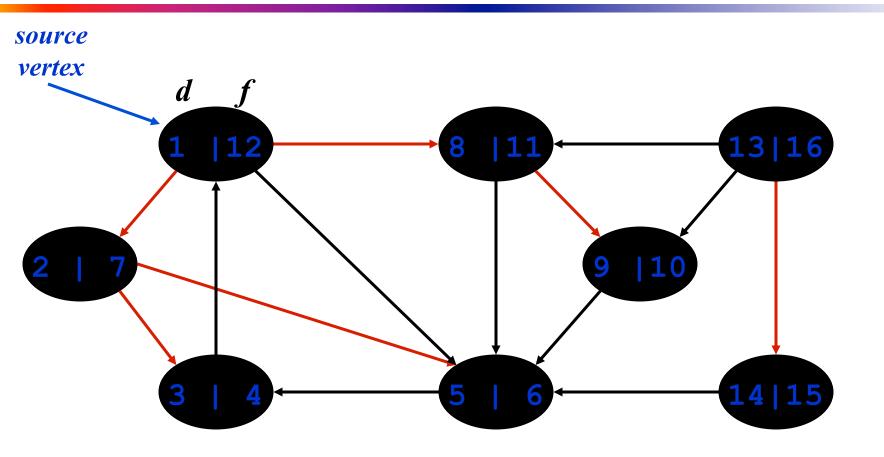






DFS: Kinds of edges

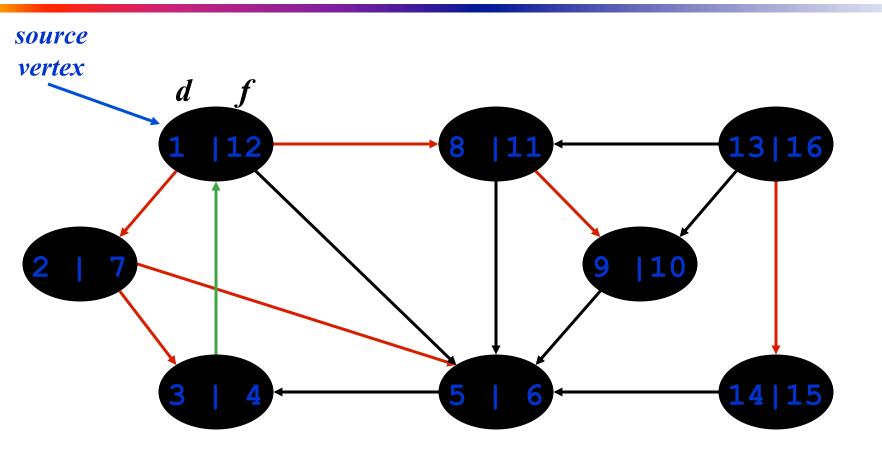
- ➤ DFS introduces an important distinction among edges in the original graph:
 - *Tree edge*: encounter new (white) vertex
 - The tree edges form a spanning forest
 - Can tree edges form cycles? Why or why not?



Tree edges

DFS: Kinds of edges

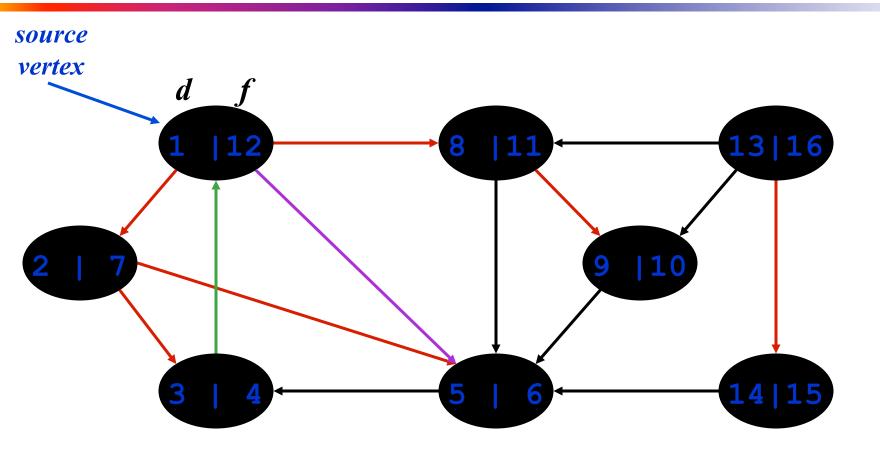
- > DFS introduces an important distinction among edges in the original graph:
 - *Tree edge*: encounter new (white) vertex
 - *Back edge*: from descendent to ancestor
 - Encounter a grey vertex (grey to grey)



Tree edges Back edges

DFS: Kinds of edges

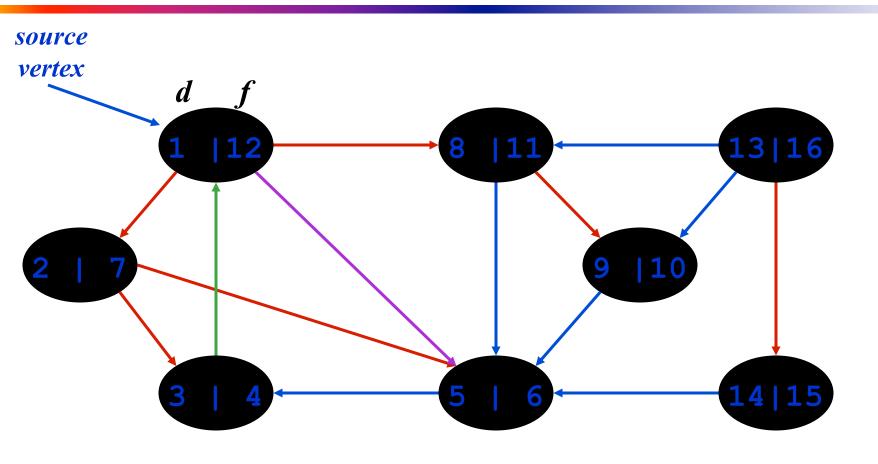
- > DFS introduces an important distinction among edges in the original graph:
 - *Tree edge*: encounter new (white) vertex
 - *Back edge*: from descendent to ancestor
 - Forward edge: from ancestor to descendent
 - Not a tree edge, though
 - From grey node to black node



Tree edges Back edges Forward edges

DFS: Kinds of edges

- > DFS introduces an important distinction among edges in the original graph:
 - *Tree edge*: encounter new (white) vertex
 - *Back edge*: from descendent to ancestor
 - Forward edge: from ancestor to descendent
 - *Cross edge*: between a tree or subtrees
 - From a grey node to a black node



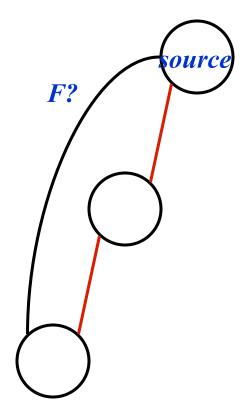
Tree edges Back edges Forward edges Cross edges

DFS: Kinds of edges

- > DFS introduces an important distinction among edges in the original graph:
 - *Tree edge*: encounter new (white) vertex
 - *Back edge*: from descendent to ancestor
 - Forward edge: from ancestor to descendent
 - *Cross edge*: between a tree or subtrees
- Note: tree & back edges are important; most algorithms don't distinguish forward & cross

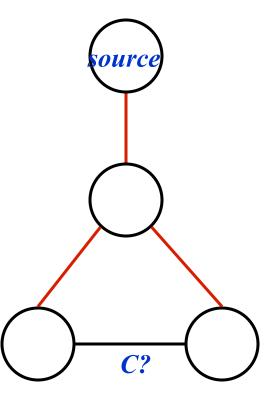
DFS: Kinds Of Edges

- Thm 23.9: If G is undirected, a DFS produces only tree and back edges
- Proof by contradiction:
 - Assume there's a forward edge
 - But F? edge must actually be a back edge (why?)



DFS: Kinds Of Edges

- Thm 23.9: If G is undirected, a DFS produces only tree and back edges
- Proof by contradiction:
 - Assume there's a cross edge
 - But C? edge cannot be cross:
 - must be explored from one of the vertices it connects, becoming a tree vertex, before other vertex is explored
 - So in fact the picture is wrong...both lower tree edges cannot in fact be tree edges



DFS And Graph Cycles

- Thm: An undirected graph is *acyclic* iff a DFS yields no back edges
 - If acyclic, no back edges (because a back edge implies a cycle
 - If no back edges, acyclic
 - No back edges implies only tree edges (*Why?*)
 - Only tree edges implies we have a tree or a forest
 - Which by definition is acyclic
- Thus, can run DFS to find whether a graph has a cycle

DFS And Cycles

➤ How would you modify the code to detect cycles?

```
DFS(G)
   for each vertex u \in G->V
      u->color = WHITE;
   time = 0;
   for each vertex u \in G->V
      if (u->color == WHITE)
         DFS Visit(u);
```

```
DFS Visit(u)
   u->color = GREY;
   time = time+1;
   u->d = time;
   for each v \in u-\lambda dj[]
       if (v->color == WHITE)
          DFS Visit(v);
   u->color = BLACK;
   time = time+1;
   u->f = time;
```

DFS And Cycles

> What will be the running time?

```
DFS(G)
   for each vertex u \in G->V
      u->color = WHITE;
   time = 0;
   for each vertex u \in G->V
      if (u->color == WHITE)
         DFS Visit(u);
```

```
DFS Visit(u)
   u->color = GREY;
   time = time+1;
   u->d = time;
   for each v \in u-\lambda dj[]
       if (v->color == WHITE)
          DFS Visit(v);
   u->color = BLACK;
   time = time+1;
   u->f = time;
```

DFS And Cycles

- > What will be the running time?
- \rightarrow A: O(V+E)

- We can actually determine if cycles exist in O(V) time:
 - In an undirected acyclic forest, $|E| \le |V| 1$
 - So count the edges: if ever see |V| distinct edges, must have seen a back edge along the way